

COMP2261 ARTIFICIAL INTELLIGENCE / MACHINE LEARNING

Multivariate Linear Regression

-- Vectorisation

Dr SHI Lei

- Vectorisation of Hypothesis Function
- Vectorisation of Cost Function
- Vectorisation of Gradient Descent

Vectorisation of Hypothesis Function

Vectorisation of Hypothesis Function

$$h_{\theta}(x) = \theta_0 + \theta_1 \cdot x_1 + \theta_2 \cdot x_2 + \cdots + \theta_n \cdot x_n$$



$$h_{\theta}(x) = \theta_0 \cdot 1 + \theta_1 \cdot x_1 + \theta_2 \cdot x_2 + \cdots + \theta_n \cdot x_n$$

$$\theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \\ \vdots \\ \theta_n \end{bmatrix} \in \mathbb{R}^{n+1}$$

parameter vector

$$x = \begin{bmatrix} 1 \\ x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} \in \mathbb{R}^{n+1}$$

variable vector

Vectorisation of Hypothesis Function

$$h_{\theta}(\mathbf{x}) = \theta_0 + \theta_1 \cdot x_1 + \theta_2 \cdot x_2 + \cdots + \theta_n \cdot x_n$$



$$h_{\theta}(\mathbf{x}) = \theta_0 \cdot x_0 + \theta_1 \cdot x_1 + \theta_2 \cdot x_2 + \cdots + \theta_n \cdot x_n \quad (x_0 = 1, \text{constant})$$

$$= [\theta_0, \theta_1, \theta_2, \dots, \theta_n] \cdot \begin{bmatrix} 1 \\ x_1 \\ x_2 \\ \dots \\ x_n \end{bmatrix}$$

$$= \boldsymbol{\theta}^T \mathbf{x} \quad \text{for single instance}$$

Vectorisation of Hypothesis Function for the whole Training Set

EXAMPLE.

Happiness predictor

features

label

income (k)	age	# of children	tea (cups/week)	happiness
38.63	46	1	7	2.314
49.79	37	0	15	3.433
40.34	52	3	20	4.03
21.18	25	0	3	1.45
...

EXAMPLE.

Happiness predictor

x_1	x_2	x_3	x_4	y
income (k)	age	# of children	tea (cups/week)	happiness
38.63	46	1	7	2.314
49.79	37	0	15	3.433
40.34	52	3	20	4.03
21.18	25	0	3	1.45
...

EXAMPLE. Happiness predictor

x_1	x_2	x_3	x_4	...	x_n	y
income (k)	age	# of children	tea (cups/week)	...	the n^{th} feature	happiness
38.63	46	1	7	...	*	2.314
49.79	37	0	15	...	*	3.433
40.34	52	3	20	...	*	4.03
21.18	25	0	3	...	*	1.45
...

EXAMPLE. Happiness predictor

		income (k)	age	# of children	tea (cups/week)	...	the n^{th} feature	happiness	
$(x^{(1)})^T$	1	38.63	46	1	7	...	*	2.314	$y^{(1)}$
$(x^{(2)})^T$	1	49.79	37	0	15	...	*	3.433	$y^{(2)}$
$(x^{(3)})^T$	1	40.34	52	3	20	...	*	4.03	$y^{(3)}$
$(x^{(4)})^T$	1	21.18	25	0	3	...	*	1.45	$y^{(4)}$
...
$(x^{(m)})^T$	1	*	*	*	*	...	*	1.45	$y^{(m)}$

EXAMPLE. Happiness predictor

$$(x^{(2)})^T = \begin{bmatrix} 1 & x_0^{(2)} & x_1^{(2)} & x_2^{(2)} & x_3^{(2)} & x_4^{(2)} & \dots & x_n^{(2)} \\ 1 & 49.79 & 37 & 0 & 15 & \dots & \star \end{bmatrix}$$

EXAMPLE.

Happiness predictor

$$x^{(2)} = \begin{bmatrix} 1 \\ 49.79 \\ 37 \\ 0 \\ 15 \\ \dots \\ * \end{bmatrix} \begin{matrix} x_0^{(2)} \\ x_1^{(2)} \\ x_2^{(2)} \\ x_3^{(2)} \\ x_4^{(2)} \\ \dots \\ x_n^{(2)} \end{matrix}$$

$$x^{(i)} = \begin{bmatrix} x_0^{(i)} \\ x_1^{(i)} \\ x_2^{(i)} \\ x_3^{(i)} \\ x_4^{(i)} \\ \dots \\ x_n^{(i)} \end{bmatrix}$$

For the i-th instance

Vectorisation of Hypothesis Function

$$X = \begin{bmatrix} (\mathbf{x}^{(1)})^T \\ (\mathbf{x}^{(2)})^T \\ (\mathbf{x}^{(3)})^T \\ (\mathbf{x}^{(4)})^T \\ \dots \\ (\mathbf{x}^{(m)})^T \end{bmatrix}$$

Vectorisation of Hypothesis Function

$$\mathbf{X} = \begin{bmatrix} (\mathbf{x}^{(1)})^T \\ (\mathbf{x}^{(2)})^T \\ (\mathbf{x}^{(3)})^T \\ (\mathbf{x}^{(4)})^T \\ \dots \\ (\mathbf{x}^{(m)})^T \end{bmatrix} = \begin{bmatrix} 1 & x_1^{(1)} & x_2^{(1)} & \dots & x_j^{(1)} & \dots & x_n^{(1)} \\ 1 & x_1^{(2)} & x_2^{(2)} & \dots & x_j^{(2)} & \dots & x_n^{(2)} \\ 1 & x_1^{(3)} & x_2^{(3)} & \dots & x_j^{(3)} & \dots & x_n^{(3)} \\ \vdots & \vdots & \vdots & \ddots & \vdots & \ddots & \vdots \\ 1 & x_1^{(i)} & x_2^{(i)} & \dots & x_j^{(i)} & \dots & x_n^{(i)} \\ \vdots & \vdots & \vdots & \ddots & \vdots & \ddots & \vdots \\ 1 & x_1^{(m)} & x_2^{(m)} & \dots & x_j^{(m)} & \dots & x_n^{(m)} \end{bmatrix} \quad \boldsymbol{\theta} = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \\ \vdots \\ \theta_j \\ \vdots \\ \theta_n \end{bmatrix}$$

Vectorisation of Hypothesis Function

$$h_{\theta} \left(x_1^{(1)}, x_2^{(1)}, \dots, x_n^{(1)} \right) = h_{\theta}(\mathbf{x}^{(1)})$$

$$\begin{aligned}
 \mathbf{X}\boldsymbol{\theta} = & \begin{bmatrix} 1 & x_1^{(1)} & x_2^{(1)} & \dots & x_j^{(1)} & \dots & x_n^{(1)} \\ 1 & x_1^{(2)} & x_2^{(2)} & \dots & x_j^{(2)} & \dots & x_n^{(2)} \\ 1 & x_1^{(3)} & x_2^{(3)} & \dots & x_j^{(3)} & \dots & x_n^{(3)} \\ \vdots & \vdots & \vdots & \ddots & \vdots & \ddots & \vdots \\ 1 & x_1^{(i)} & x_2^{(i)} & \dots & x_j^{(i)} & \dots & x_n^{(i)} \\ \vdots & \vdots & \vdots & \ddots & \vdots & \ddots & \vdots \\ 1 & x_1^{(m)} & x_2^{(m)} & \dots & x_2^{(m)} & \dots & x_n^{(m)} \end{bmatrix} \cdot \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \\ \vdots \\ \theta_j \\ \vdots \\ \theta_n \end{bmatrix} = \begin{bmatrix} \theta_0 + \theta_1 \cdot x_1^{(1)} + \theta_2 \cdot x_2^{(1)} + \dots + \theta_n \cdot x_n^{(1)} \\ \theta_0 + \theta_1 \cdot x_1^{(2)} + \theta_2 \cdot x_2^{(2)} + \dots + \theta_n \cdot x_n^{(2)} \\ \theta_0 + \theta_1 \cdot x_1^{(3)} + \theta_2 \cdot x_2^{(3)} + \dots + \theta_n \cdot x_n^{(3)} \\ \dots \\ \theta_0 + \theta_1 \cdot x_1^{(i)} + \theta_2 \cdot x_2^{(i)} + \dots + \theta_n \cdot x_n^{(i)} \\ \dots \\ \theta_0 + \theta_1 \cdot x_1^{(m)} + \theta_2 \cdot x_2^{(m)} + \dots + \theta_n \cdot x_n^{(m)} \end{bmatrix} \\
 & \mathbb{R}^{m \times (n+1)} \qquad \qquad \mathbb{R}^{(n+1) \times 1} \qquad \qquad \mathbb{R}^{m \times 1}
 \end{aligned}$$

Vectorisation of Hypothesis Function

$$\begin{aligned}
 \mathbf{X}\boldsymbol{\theta} &= \begin{bmatrix} 1 & x_1^{(1)} & x_2^{(1)} & \cdots & x_j^{(1)} & \cdots & x_n^{(1)} \\ 1 & x_1^{(2)} & x_2^{(2)} & \cdots & x_j^{(2)} & \cdots & x_n^{(2)} \\ 1 & x_1^{(3)} & x_2^{(3)} & \cdots & x_j^{(3)} & \cdots & x_n^{(3)} \\ \vdots & \vdots & \vdots & \ddots & \vdots & \ddots & \vdots \\ 1 & x_1^{(i)} & x_2^{(i)} & \cdots & x_j^{(i)} & \cdots & x_n^{(i)} \\ \vdots & \vdots & \vdots & \ddots & \vdots & \ddots & \vdots \\ 1 & x_1^{(m)} & x_2^{(m)} & \cdots & x_j^{(m)} & \cdots & x_n^{(m)} \end{bmatrix} \cdot \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \\ \vdots \\ \theta_j \\ \vdots \\ \theta_n \end{bmatrix} = \begin{bmatrix} h_{\boldsymbol{\theta}}(\mathbf{x}^{(1)}) \\ h_{\boldsymbol{\theta}}(\mathbf{x}^{(2)}) \\ h_{\boldsymbol{\theta}}(\mathbf{x}^{(3)}) \\ \cdots \\ h_{\boldsymbol{\theta}}(\mathbf{x}^{(i)}) \\ \cdots \\ h_{\boldsymbol{\theta}}(\mathbf{x}^{(m)}) \end{bmatrix} \\
 &\quad \mathbb{R}^{m \times (n+1)} \qquad \qquad \mathbb{R}^{(n+1) \times 1} \qquad \qquad \mathbb{R}^{m \times 1}
 \end{aligned}$$

Vectorisation of Hypothesis Function

$$\mathbf{X}\boldsymbol{\theta} = \begin{bmatrix} 1 & x_1^{(1)} & x_2^{(1)} & \dots & x_j^{(1)} & \dots & x_n^{(1)} \\ 1 & x_1^{(2)} & x_2^{(2)} & \dots & x_j^{(2)} & \dots & x_n^{(2)} \\ 1 & x_1^{(3)} & x_2^{(3)} & \dots & x_j^{(3)} & \dots & x_n^{(3)} \\ \vdots & \vdots & \vdots & \ddots & \vdots & \ddots & \vdots \\ 1 & x_1^{(i)} & x_2^{(i)} & \dots & x_j^{(i)} & \dots & x_n^{(i)} \\ \vdots & \vdots & \vdots & \ddots & \vdots & \ddots & \vdots \\ 1 & x_1^{(m)} & x_2^{(m)} & \dots & x_j^{(m)} & \dots & x_n^{(m)} \end{bmatrix} \cdot \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \\ \vdots \\ \theta_j \\ \vdots \\ \theta_n \end{bmatrix} = \begin{bmatrix} h_{\boldsymbol{\theta}}(\mathbf{x}^{(1)}) \\ h_{\boldsymbol{\theta}}(\mathbf{x}^{(2)}) \\ h_{\boldsymbol{\theta}}(\mathbf{x}^{(3)}) \\ \dots \\ h_{\boldsymbol{\theta}}(\mathbf{x}^{(i)}) \\ \dots \\ h_{\boldsymbol{\theta}}(\mathbf{x}^{(m)}) \end{bmatrix}$$

Vectorised Hypothesis Function

$\longrightarrow \mathbf{h}_{\boldsymbol{\theta}}(\mathbf{x}) = \mathbf{X} \boldsymbol{\theta}$

vector / matrix

A vector containing the results of the hypothesis function, i.e., the predicted value, for all the instances from the training set.

Vectorisation of Hypothesis Function

$$h_{\theta}(x) = \theta^T x$$

vs

$$h_{\theta}(x) = X \theta$$

takes single instance

outputs predicted label

takes all instances

outputs a vector of predicted labels

Vectorisation of Cost Function

Vectorisation of Cost Function

$$\mathbf{X}\boldsymbol{\theta} - \mathbf{y} = \begin{bmatrix} 1 & x_1^{(1)} & x_2^{(1)} & \cdots & x_n^{(1)} \\ 1 & x_1^{(2)} & x_2^{(2)} & \cdots & x_n^{(2)} \\ 1 & x_1^{(3)} & x_2^{(3)} & \cdots & x_n^{(3)} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_1^{(m)} & x_2^{(m)} & \cdots & x_n^{(m)} \end{bmatrix} \cdot \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \\ \vdots \\ \theta_n \end{bmatrix} - \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ y^{(3)} \\ \vdots \\ y^{(m)} \end{bmatrix} = \begin{bmatrix} h_{\boldsymbol{\theta}}(x^{(1)}) - y^{(1)} \\ h_{\boldsymbol{\theta}}(x^{(2)}) - y^{(2)} \\ h_{\boldsymbol{\theta}}(x^{(3)}) - y^{(3)} \\ \vdots \\ h_{\boldsymbol{\theta}}(x^{(m)}) - y^{(m)} \end{bmatrix}$$

Vectorisation of Cost Function

$$\mathbf{X}\boldsymbol{\theta} - \mathbf{y} = \begin{bmatrix} h_{\theta}(x^{(1)}) - y^{(1)} \\ h_{\theta}(x^{(2)}) - y^{(2)} \\ h_{\theta}(x^{(3)}) - y^{(3)} \\ \vdots \\ h_{\theta}(x^{(m)}) - y^{(m)} \end{bmatrix}$$

Vectorisation of Cost Function

$$\begin{aligned} & (\mathbf{X}\boldsymbol{\theta} - \mathbf{y})^T \cdot (\mathbf{X}\boldsymbol{\theta} - \mathbf{y}) \\ &= \begin{bmatrix} h_{\theta}(x^{(1)}) - y^{(1)}, h_{\theta}(x^{(2)}) - y^{(2)}, h_{\theta}(x^{(3)}) - y^{(3)}, \dots, h_{\theta}(x^{(m)}) - y^{(m)} \end{bmatrix} \cdot \begin{bmatrix} h_{\theta}(x^{(1)}) - y^{(1)} \\ h_{\theta}(x^{(2)}) - y^{(2)} \\ h_{\theta}(x^{(3)}) - y^{(3)} \\ \vdots \\ h_{\theta}(x^{(m)}) - y^{(m)} \end{bmatrix} \\ &= (h_{\theta}(x^{(1)}) - y^{(1)})^2 + (h_{\theta}(x^{(2)}) - y^{(2)})^2 + (h_{\theta}(x^{(3)}) - y^{(3)})^2 + \dots + (h_{\theta}(x^{(m)}) - y^{(m)})^2 \\ &= \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 \end{aligned}$$

Vectorisation of Cost Function



$$(\mathbf{X}\boldsymbol{\theta} - \mathbf{y})^T \cdot (\mathbf{X}\boldsymbol{\theta} - \mathbf{y}) = \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Vectorisation of Cost Function

$$\sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 = (\mathbf{X}\boldsymbol{\theta} - \mathbf{y})^T (\mathbf{X}\boldsymbol{\theta} - \mathbf{y})$$

Vectorisation of Cost Function

$$\frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 = \frac{1}{2m} (\mathbf{X}\boldsymbol{\theta} - \mathbf{y})^T (\mathbf{X}\boldsymbol{\theta} - \mathbf{y})$$

Vectorisation of Cost Function

$$J(\boldsymbol{\theta}) = \frac{1}{2m} \sum_{i=1}^m (h_{\boldsymbol{\theta}}(x^{(i)}) - y^{(i)})^2 = \frac{1}{2m} (\mathbf{X}\boldsymbol{\theta} - \mathbf{y})^T (\mathbf{X}\boldsymbol{\theta} - \mathbf{y})$$

Vectorisation of Cost Function

$$J(\boldsymbol{\theta}) = \frac{1}{2m} (\mathbf{X}\boldsymbol{\theta} - \mathbf{y})^T (\mathbf{X}\boldsymbol{\theta} - \mathbf{y})$$

Vectorisation of Gradient Descent

Vectorisation of Gradient Descent

Repeat until convergence {

$$\theta_0 := \theta_0 - \alpha \frac{\partial J}{\partial \theta_0}(\boldsymbol{\theta})$$

$$\theta_1 := \theta_1 - \alpha \frac{\partial J}{\partial \theta_1}(\boldsymbol{\theta})$$

...

$$\theta_n := \theta_n - \alpha \frac{\partial J}{\partial \theta_n}(\boldsymbol{\theta})$$

}

$$\begin{aligned} \boxed{\frac{\partial J}{\partial \theta_k}(\boldsymbol{\theta})} &= \frac{1}{m} \sum_{i=1}^m \boxed{(h_{\theta}(x)^{(i)} - y^{(i)}) x_k^{(i)}} \\ &= \frac{1}{m} \sum_{i=1}^m x_k^{(i)} (h_{\theta}(x)^{(i)} - y^{(i)}) \\ &= \frac{1}{m} \mathbf{x}_k^T (\mathbf{X}\boldsymbol{\theta} - \mathbf{y}) \end{aligned}$$

Vectorisation of Gradient Descent

Repeat until convergence {

$$\begin{aligned} \theta_0 &:= \theta_0 - \alpha \frac{\partial J}{\partial \theta_0}(\boldsymbol{\theta}) \\ \theta_1 &:= \theta_1 - \alpha \frac{\partial J}{\partial \theta_1}(\boldsymbol{\theta}) \\ &\dots \\ \theta_n &:= \theta_n - \alpha \frac{\partial J}{\partial \theta_n}(\boldsymbol{\theta}) \end{aligned}$$

$\frac{\partial J}{\partial \theta_k}(\boldsymbol{\theta}) = \frac{1}{m} \mathbf{x}_k^T (\mathbf{X}\boldsymbol{\theta} - \mathbf{y})$

}

Vectorisation of Gradient Descent

Repeat until convergence {

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \mathbf{x}_0^T (\mathbf{X}\boldsymbol{\theta} - \mathbf{y})$$

$$\theta_1 := \theta_1 - \alpha \frac{1}{m} \mathbf{x}_1^T (\mathbf{X}\boldsymbol{\theta} - \mathbf{y})$$

...

$$\theta_n := \theta_n - \alpha \frac{1}{m} \mathbf{x}_n^T (\mathbf{X}\boldsymbol{\theta} - \mathbf{y})$$

}

Vectorisation of Gradient Descent

Repeat until convergence {

$$\boldsymbol{\theta} := \boldsymbol{\theta} - \alpha \frac{1}{m} \mathbf{X}^T (\mathbf{X}\boldsymbol{\theta} - \mathbf{y})$$

}

- Vectorisation of Hypothesis Function

$$h_{\theta}(x) = \theta^T x \quad \mathbf{h}_{\theta}(x) = X \theta$$

- Vectorisation of Cost Function

$$J(\theta) = \frac{1}{2m} (\mathbf{X}\theta - \mathbf{y})^T (\mathbf{X}\theta - \mathbf{y})$$

- Vectorisation of Gradient Descent

Repeat until convergence {

$$\theta := \theta - \alpha \frac{1}{m} X^T (\mathbf{X}\theta - \mathbf{y})$$

}

Comparing with un-vectorised

- ✓ Easier to implement
- ✓ Computationally more efficient